

Biological Stats II : Lab 1

Gavin Fay

01/18/2023

Lab schedule

1/18: Introduction to R and R Studio, working with data

1/25: Intro to Visualization

2/01: Probability, linear modeling

2/08: Data wrangling, model summaries

2/15: Iteration

2/22: Creating functions, debugging

3/01: Simulation, Resampling

3/15: Flex: more modeling (brms, glmmTMB)

3/29: Spatial data or tidymodeling

Format for lab component

Description/overview of concepts by instructor

Live coding. Interactive, follow along yourselves

Lab exercises to practice material

Assignment each week is to finish the exercises

Extra credit: #TidyTuesday

Community support

“Try stuff and see what happens”

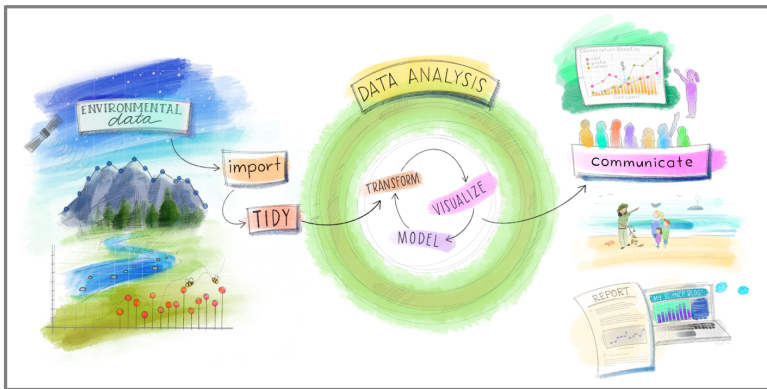
Why ?

- Reproducible Science!
- Full data science workflow
- Language & environment for statistical computing & graphic
- Open source & Free
- Lots of scientists use it!
- AMAZING online community
- Works well with other tools



Artwork by [@allison_horst](https://twitter.com/allison_horst)

Why ?



Why ?

- Reproducible Science!
- Full data science workflow
- Language & environment for statistical computing & graphic
- Open source & Free
- Lots of scientists use it!
- AMAZING online community
- Works well with other tools



Artwork by [@allison_horst](https://twitter.com/allison_horst)

Why R?

Reproducible

- command line interface encourages organization
- scripts allow others (and you!) to reproduce analyses from end-to-end
- integration of analysis with document creation

Extensible

- new methods delivered as developed
- continual expansion through new packages

Open-source

- all code can be examined by the user

Free

- available to large set of users (and therefore developers)

R is not the only solution out there.

The real goal is not to teach R, but concepts that all programming depends on.


Trevor Branch rule:

*“Every analysis you do on a dataset will have to be redone
10–15 times before publication. Plan accordingly.”*

Recommended reading

R for Data Science, 2017 (Wickham and Grolemund)

<https://r4ds.had.co.nz/>



Recommended reading

An introduction to R (Venables et al.)

- <http://cran.r-project.org/doc/manuals/R-intro.pdf>

R reference card 2.0 (Baggott)

- <http://cran.r-project.org/doc/contrib/Baggott-refcard-v2.pdf>

- Extremely useful handout: put on wall in view of your desk

Cheatsheets

- <https://posit.co/resources/cheatsheets/>

There are many (many) R books out there. Good for reference. e.g.

- The R Book, 2007 (Crawley)

- Modern Applied Statistics with S, 20002 (Venables and Ripley)

- Introductory Statistics with R, 2002 (Dalgaard)

Installing R and R Studio

R (<http://r-project.org>)

- download appropriate version for your OS

R Studio (<https://posit.co/download/rstudio-desktop/>)

- a very good Integrated Development Environment (IDE) for R provides:
 - text editor
 - syntax highlighting
 - seamless code execution with R

You can use other text editors with R, but RStudio is well organized.
(also looks the same regardless of Operating System)

Using R today: RStudio Cloud

An alternative:

- ▶ web browser interface

<https://rstudio.cloud/spaces/47288/project/856390>

Getting started

Enter instructions at R console command line prompt (>):

e.g. type

```
> 2 + 2
```

R acts as a calculator and returns (prints) the result.

```
> 2 + 2  
[1] 4
```

Simple commands

```
> 3^2
[1] 9
> 2*(2+2)
[1] 8
> 2*2+2
[1] 6
> log(10)
[1] 2.302585
> exp(1)
[1] 2.718282
> x <- 3
> 2*x
[1] 6
```

The `<-` means 'assign'. i.e. 'assign a value of 3 to the variable x'.
`<-` is preferable to using `=`

Scripts and RStudio

- ▶ Typing commands into the console can get tedious.
- ▶ Scripts are text files containing lines of code.
- ▶ Scripts provide a complete record of analyses.
- ▶ Code can be run (executed) from these files repeatedly.
- ▶ Scripts can be created in a text editor and copied into the R console.

Or...

- ▶ RStudio integrates scripts, R console, and output in a user-friendly development environment.
- ▶ To run code in RStudio, select code and type
Ctrl+Enter (Windows)
Command+Enter (Mac)
The code will run in the R console.

RStudio

The image shows the RStudio application window. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Tools, Window, and Help. The title bar indicates the path ~/classes/biostat2/labs - master - RStudio and the user Gavin Fay. The interface is divided into four main panes:

- Source Editor (Top Left):** Contains an R script file named lab1.R with the following code:

```
1 # Lab 1 code
2
3 Z<2
4 x <- 1:10
5 y <- rnorm(10)
6 plot(x~y,pch=16,ylim=c(-2,2))
7 abline(h=0,col=gray(0.2))
8
9
```

An annotation box labeled "Files with R code (scripts)" points to this pane.
- Environment (Top Right):** Shows the "Global Environment" with the message "Environment is empty". An annotation box labeled "Details of objects you have created." points to this pane.
- Console (Bottom Left):** Displays the R startup message and the results of running the code in the source editor:

```
> Z<2
[1] 4
>
```

An annotation box labeled "R Console, command prompt. Results of running code" points to this pane.
- Help Viewer (Bottom Right):** Displays the documentation for "Generic X-Y Plotting". The title is "Generic X-Y Plotting" and the description is "Generic function for plotting". The usage is `plot(x, y, ...)`. The arguments section lists: `x` (the coordinates of points in the plot), `y` (the y coordinates of points in the plot), and `...` (arguments to be passed to methods). An annotation box labeled "Help documentation, Plots" points to this pane.

Projects in RStudio

Good to keep all the files associated with a particular project organized together – input data, R scripts, analytical results, figures.

RStudio has built-in support for this via its projects.

Projects provide an easy organization of your files and analyses.

R runs from a working directory, the location (folder) on your computer that R will use as a default when looking for data files, producing output, etc.

Project files (`.Rproj`)

When you open a project, RStudio sets the working directory to the folder in which the `.Rproj` file exists.

Using projects can be VERY helpful.

Getting Help in R


R is a programming language, there is a learning curve.

Fortunately, there are lots of resources:

NOT XKCD by Kai



PRIORITIES



(FORGET WHAT HE ASKED FOR. CAN WE TALK ABOUT THE FACT YOU SUMMONED A WISH GRANTING COMPUTER DEMON?)

Don't summon a wish-granting computer demon.

Getting Help in R

R is a programming language, there is a learning curve.

Fortunately, there are lots of resources:

- help files
- online search results
- books
- colleagues
- RStudio Community

```
> ?(mean)
> help("mean")
```

The above both get help for the function mean.

Use `help.search("function.name")` to search across packages.

`str(object.name)` shows the structure of an object.

R Help files (`?mean`)

Common format:

- ▶ Description (what the function does)
- ▶ Usage (how to use it)
- ▶ Arguments (what the function needs, options)
- ▶ Value (what does the function return)
- ▶ See Also (related functions)
- ▶ Examples (sample code showing how the function works)

Read function documentation and explore behavior by running examples!

Lab exercise 1/4

(Instructions also in lab-exercise-01.pdf)

Start a new project. (call it biostats2 or something similar)

Open a new R markdown file. Save it. (name it lastname_lab1.Rmd or something similar)

At the top of the script, add comments with your name and lab 1.

Work in pairs or individually.

Submit your .Rmd file via myCourses before lab next week.

Write code that evaluates the following when run.

$$7 + 5(4 + 3)$$

$$e^{-5(0.2+0.15)}$$

$$\frac{\sqrt{1 + 2(3 + 2)}}{\ln(3^2 + 2)}$$

Objects

Common types of objects

- Numbers
- Characters (i.e. text or strings)
- Tables
- Vectors and matrices
- Plots
- Statistical output
- Functions

Objects in R are global

Viewing objects: In RStudio see top-right Workspace tab

More generally:

```
> print(myobject)
> myobject
```

`ls()` lists all objects in the workspace.

Use `rm()` to remove an object.

Data types (modes)

Describe how objects are stored in computer memory.

In R you do ***not*** need to specify the data type.

Common data types:

- ▶ Numeric (integer, floating point numbers or doubles)
- ▶ Logical (Boolean, true or false)
- ▶ Characters (text or string data)

Types are not always obvious in R, but can be important to know.

Data types II

```
> myobject <- log(10)
> mode(myobject)
[1] "numeric"
> is.numeric(myobject)
[1] TRUE
> typeof(myobject)
[1] "double"
> newobject <- as.integer(myobject)
> typeof(newobject)
[1] "integer"
> is.character(myobject)
[1] FALSE
> typeof("hello world")
[1] "character"
```


Vectors

```
> weights <- c(2.3,5.4,7.5,9)
> print(weights)
[1] 2.3 5.4 7.5 9.0
> years <- 2007:2016
> print(years)
[1] 2007 2008 2009 2010 2011 2012 2013 2014 2015 2016
> years <- seq(from=2000,to=2016,by=2)
> print(years)
[1] 2000 2002 2004 2006 2008 2010 2012 2014 2016
> x <- rep(3,times=10)
> print(x)
[1] 3 3 3 3 3 3 3 3 3 3
> rep(1:3,times=3)
[1] 1 2 3 1 2 3 1 2 3
> rep(1:3,length=10)
[1] 1 2 3 1 2 3 1 2 3 1
```

More on Vectors

Vectors are ordered and can be referred to by element(s) using

```
> (years <- 2007:2016)
[1] 2007 2008 2009 2010 2011 2012 2013 2014 2015 2016
> years[3]
[1] 2009
> years[5:6]
[1] 2011 2012
> which(years==2010)
[1] 4
> years[-c(2,4)] # A negative index excludes elements
[1] 2007 2009 2011 2012 2013 2014 2015 2016
```

Vector operations are element-wise

```
> (x <- 1:5)
[1] 1 2 3 4 5
> 2*x
```

Lab exercise 2/4

(Instructions also in lab-exercise-01.pdf)

Create vectors using `seq()`, `rep()`, and mathematical operators.
Only use `c()` when absolutely necessary.

hint Remember you can get help on a function by typing
`?functionname`

- ▶ Positive integers from 1 to 99
- ▶ Odd integers between 1 and 99
- ▶ The numbers 1,1,1, 2,2,2, 3,3,3
- ▶ The numbers -5,-4,-3,-5,-4,-3,-5,-4,-3
- ▶ The fractions 1, 1/2, 1/3, 1/4, ..., 1/10
- ▶ The cubes 1, 8, 27, 64, 125, 216

Useful functions

```
> x <- c(5,3,2,6,3,9,1,18)
> length(x)  # length of vector x
[1] 8
> sort(unique(x))  # sorted vector of unique values in x
[1] 1 2 3 5 6 9 18
> min(x)        # minimum value in x
[1] 1
> max(x)        # maximum value in x
[1] 18
> mean(x)       # mean of x
[1] 5.875
> median(x)     # median of x
[1] 4
> sd(x)         # standard deviation of x
[1] 5.514591
> range(x)      # range of values in x
[1] 1 18
> range(x)[2]   # 2nd element of values returned by range()
[1] 18
> quantile(x)  # optional argument 'probs' can be handy
  0%   25%   50%   75%  100%
1.00  2.75  4.00  6.75 18.00
```

Boolean logic operators

Operator	R Code
AND	& (&&)
OR	()
NOT	!
less than	<
greater than	>
less than or equal	<=
greater than or equal	>=
equals	==
NOT equal	!=

&& and || are used when asking IF statements.
These only use a single value, not a vector.

Boolean examples

```
> x <- 7
> x == 7
[1] TRUE
> x < 10
[1] TRUE
> x < -3
[1] FALSE
> x > 0 & x <= 12
[1] TRUE
> x >= 10 | x < 0
[1] FALSE
```

```
> y <- c(4,8)
> y > 5 #returns a logical vector
[1] FALSE TRUE
> y[y>5] #returns elements of y that meet condition
[1] 8
> which(y>5) #index of y that meets condition
[1] 2
> any(y>5)
[1] TRUE
> all(y>5)
[1] FALSE
```

Lab exercise 3/4

Complete the following using the vector `y`:

```
y <- c(3,2,15,-1,22,1,9,17,5)
```

- ▶ Display the first and last values.
- ▶ Find the last value for a vector of any length.
- ▶ Display the values that are greater than the mean of `y`.
- ▶ Display the positions (indices) of the values greater than the mean.
- ▶ Are all the values positive?
- ▶ Are any of the values equal to the mean?
- ▶ Are any of the values equal to the median?

Other types of objects

matrices (more generally, arrays)

- multi-dimensional generalizations of vectors.
- are vectors that can be indexed by two or more indices.

factors

- compact ways to handle categorical data.

lists

- general form of vector, elements need not be the same type.
- elements often themselves vectors or lists.
- convenient way to return results of statistical computations.

dataframes

- matrix-like structures, columns can be of different types.
- often 'data matrices' with one row per observational unit but with (possibly) both numerical and categorical variables.
- experiments are often best described by data frames: treatments are categorical but the response is numeric.

functions

- are themselves objects in R which can be stored in the project's workspace.
- provide a simple and convenient way to extend R.

Dataframes

There are lots of data set examples in R.
e.g. record times for 35 Scottish hill races

```
> library(MASS)
> head(hills,n=3) # shows first few lines. Also tail()
      dist climb  time
Greenmantle  2.5   650 16.083
Carnethy     6.0  2500 48.350
Craig Dunain  6.0   900 33.650
> names(hills) # get the names of the data frame
[1] "dist" "climb" "time"
```

Creating dataframes

```
> fish <- c("cod","haddock","dogfish","pollock")
> length <- c(34,23,75,18)
> age <- c(6,3,17,2)
> fish.data <- data.frame(fish=fish,length=length,age=age)
> head(fish.data)
  fish length age
1   cod     34  6
2 haddock    23  3
3 dogfish    75 17
4 pollock    18  2
```

Extracting information from data frames

Use the \$ to extract vectors from a data frame

```
> hills$dist
[1] 2.5 6.0 6.0 7.5 8.0 8.0 16.0 6.0 5.0 6.0 28.0 5.0
[14] 6.0 4.5 10.0 14.0 3.0 4.5 5.5 3.0 3.5 6.0 2.0 3.0
[27] 6.0 5.0 6.5 5.0 10.0 6.0 18.0 4.5 20.0
```

You can also specify the row index, column index, or both
object[row,column]

```
> # extract the element in row 1, column 2
> hills[1,2]
[1] 650
> hills$climb[1]
[1] 650
> # extract the first row
> hills[1,]
      dist climb  time
Greenmantle 2.5  650 16.083
```

```

> # extract all of column 2
> hills[,2]      # also hills[, "climb"]
[1] 650 2500 900 800 3070 2866 7500 800 800 650 2100 2000
[14] 500 1500 3000 2200 350 1000 600 300 1500 2200 900 600
[27] 800 950 1750 500 4400 600 5200 850 5000
> # exclude column 1, but retain the other columns (1st 3 rows)
> hills[1:3,-1]
      climb  time
Greenmantle 650 16.083
Carnethy    2500 48.350
Craig Dunain 900 33.650
> # extract rows 4 and 7
> hills[c(4,7),]
      dist climb  time
Ben Rha      7.5 800 45.600
Bens of Jura 16.0 7500 204.617
> # extract the rows that are specified by the object x
> x <- c(4,7,nrow(hills))
> hills[x,]
      dist climb  time
Ben Rha      7.5 800 45.600
Bens of Jura 16.0 7500 204.617
Moffat Chase 20.0 5000 159.833

```

Extracting information from data frames (2)

The 'tidyverse' is a popular set of R packages for interacting with data and performing statistical analyses.

Data frames are commonly used in a structure called a 'tibble'

```
> #install.packages('tidyverse')  
> library(tidyverse)
```

```
> hills_tbl <- as_tibble(hills)  
> hills_tbl
```

```
# A tibble: 35 x 3
```

	dist	climb	time
	<dbl>	<int>	<dbl>
1	2.5	650	16.1
2	6	2500	48.4
3	6	900	33.6
4	7.5	800	45.6
5	8	3070	62.3
6	8	2866	73.2
7	16	7500	205.
8	6	800	36.4
9	5	800	29.8
10	6	650	39.8

```
# ... with 25 more rows
```

Extract variables using the select verb function in the dplyr package

```
> dplyr::select(hills,dist)
```

	dist
--	------

Greenmantle	2.5
-------------	-----

Carnethy	6.0
----------	-----

Craig Dunain	6.0
--------------	-----

Ben Rha	7.5
---------	-----

Ben Lomond	8.0
------------	-----

Goatfell	8.0
----------	-----

Bens of Jura	16.0
--------------	------

Cairnpapple	6.0
-------------	-----

Scolty	5.0
--------	-----

Traprain	6.0
----------	-----

Lairig Ghru	28.0
-------------	------

Dollar	5.0
--------	-----

Lomonds	9.5
---------	-----

Cairn Table	6.0
-------------	-----

Eildon Two	4.5
------------	-----

Cairngorm	10.0
-----------	------

Seven Hills	14.0
-------------	------

Knock Hill	3.0
------------	-----

Black Hill	4.5
------------	-----

Creag Beag	5.5
------------	-----

Kildcon Hill	3.0
--------------	-----

...	...
-----	-----

Create a new variable of the race names

```
> hills_tbl <- mutate(hills_tbl, race=rownames(hills))
> hills_tbl <- dplyr::select(hills_tbl, race, everything())
> hills_tbl
```


A tibble: 35 x 4

	race <chr>	dist <dbl>	climb <int>	time <dbl>
1	Greenmantle	2.5	650	16.1
2	Carnethy	6	2500	48.4
3	Craig Dunain	6	900	33.6
4	Ben Rha	7.5	800	45.6
5	Ben Lomond	8	3070	62.3
6	Goatfell	8	2866	73.2
7	Bens of Jura	16	7500	205.
8	Cairnpapple	6	800	36.4
9	Scolty	5	800	29.8
10	Traprain	6	650	39.8

... with 25 more rows

Quickly visualize the dataset

```
library(ggplot2)
ggplot(data = hills_tbl, mapping = aes(x=dist,y=time)) +
  geom_point(aes(size = climb), alpha = 1/3) +
  xlab("Distance (miles)") +
  ylab("Record Time (mins)") +
  xlim(0,30) + ylim(0,225)
```



Extracting elements logically

```
> fish <- c("cod","haddock","dogfish","pollock")
> length <- c(34,23,75,18)
> age <- c(6,3,17,2)
> fish_data <- tibble(fish=fish,length=length,age=age)
> fish_data$age      # a vector
[1]  6  3 17  2
> fish_data$age > 5   # a logical vector
[1] TRUE FALSE TRUE FALSE
> filter(fish_data, age>5) #avoids having to create a logical vector
# A tibble: 2 x 3
  fish    length  age
<chr>    <dbl> <dbl>
1 cod         34     6
2 dogfish     75    17
> # combining conditions
> filter(fish_data, age > 5, fish == "dogfish")
# A tibble: 1 x 3
  fish    length  age
<chr>    <dbl> <dbl>
1 dogfish     75    17
> filter(fish_data, length < 25 | length >= 50)
# A tibble: 3 x 3
  fish    length  age
<chr>    <dbl> <dbl>
1 haddock     23     3
```


Using the Pipe

Combine verb operations with the pipe operator (`|>`) e.g.

```
> fish_data |>
+   mutate(weight = (length/100)^3) |>
+   filter(age > 5)
# A tibble: 2 x 4
  fish    length  age weight
<chr>    <dbl> <dbl>  <dbl>
1 cod         34     6  0.0393
2 dogfish     75    17  0.422
```

The pipe passes the result from one function call to the next.

Can be read as “and then”.

(also the older tidyverse pipe: `%>%`)

Summarizing Data

After selecting data we might like to perform analyses on it.

We can use `group_by()` to recognize structure in the data.

Then apply functions to variables using `summarize()`

Recall `gapminder` data, find the average life expectancy in 2007 for each continent

```
> #install.packages('gapminder')
> library(gapminder)
> gapminder |>
+   filter(year == 2007) |>
+   group_by(continent) |>
+   summarize(avg_lifeExp = mean(lifeExp))
# A tibble: 5 x 2
  continent avg_lifeExp
  <fct>      <dbl>
1 Africa      54.8
2 Americas    73.6
3 Asia        70.7
4 Europe      77.6
5 Oceania     80.7
```

Tips and Tricks

Comments

Use comments within R code to document the purpose of your code. Anything on a line after a # is ignored by R. RStudio uses a different color to help readability.

SAVE your scripts(!), not workspaces. Use meaningful variable names. Adopt a coding style and use consistently.
(e.g. <https://google.github.io/styleguide/Rguide.xml>)

The `str()` function can be incredibly helpful when querying objects.

```
> str(hills)
'data.frame':   35 obs. of  3 variables:
 $ dist : num  2.5 6 6 7.5 8 8 16 6 5 6 ...
 $ climb: int  650 2500 900 800 3070 2866 7500 800 800 650 ...
 $ time : num  16.1 48.4 33.6 45.6 62.3 ...
```

Make use of `help()` documentation.

There are almost always multiple ways of getting the same result. We'll mostly use low level functions to help you understand how R works. Some advanced functions are cleaner and do things more quickly.

Lab exercise 4/4 (data frames using hills)

1. Display the first 5 rows of the `hills` dataframe.
2. Find the fastest time.
3. Display the hill races (and distance, climbs, and times) with the 3 fastest times.
4. Extract and display the record time for Cairngorm.
5. Find how many hill races have a climb greater than the mean.
6. Display the names of the hill races that have a climb greater than the mean.
7. Display the names and times of the races that are at least 10 miles long and have a climb greater than 4000 feet.
8. Find the positions (indices) of hills that either have a climb greater than 5000 feet or have a record time less than 20 minutes.
9. Find the standard deviation of the record times for all races except for the highest climb, the Bens of Jura.
10. Display the range (minimum and maximum) of the average speed for the races.
11. Find the race that had the fastest average speed.
12. **BONUS** Find the mean of the record times for races whose names start with letters A through K.

Next time...

1/18: Introduction to R and R Studio, working with data

1/25: Intro to Visualization

2/01: Probability, linear modeling

2/08: Data wrangling, model summaries

2/15: Iteration

2/22: Creating functions, debugging

3/01: Simulation, Resampling

3/15: Flex: more modeling (brms, glmmTMB)

3/29: Spatial data or tidymodeling