

Bio Stats II : Lab 3

Acknowledgements: Punt & Branch Labs (U Washington)

Gavin Fay

01/31/2023

Lab schedule

1/18: Introduction to R and R Studio, working with data

1/25: Intro to Visualization

1/31: Probability, linear modeling

2/08: Data wrangling, model summaries

2/15: Iteration

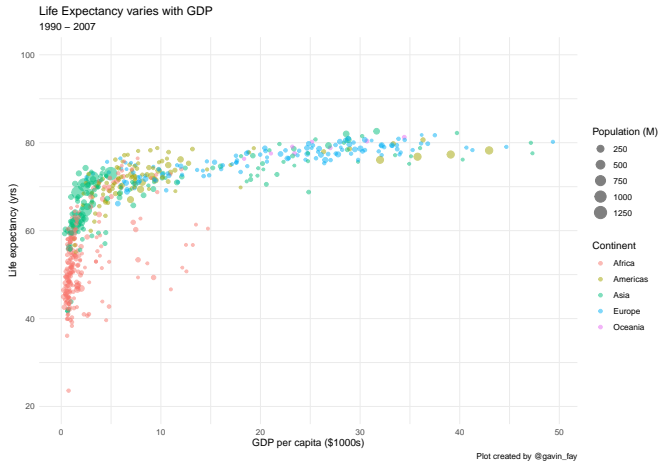
2/22: Creating functions, debugging

3/01: Simulation, Resampling

3/15: Flex: more modeling (brms, glmmTMB)

3/29: Spatial data or tidymodeling

Review



```
> suppressPackageStartupMessages(library(tidyverse))
> library(gapminder)
> gapdata <- gapminder %>%
+   filter(year >= 1990 & year <= 2007) %>%
+   group_by(continent)
> ggplot(gapdata, aes(x = gdpPercap/1000,
+                     y = lifeExp,
+                     size = pop/1000000,
+                     color = continent)) +
+   geom_point(alpha = 0.5) +
+   ylim(20,100) +
+   labs(title = "Life Expectancy varies with GDP",
+        subtitle = "1990 - 2007",
+        x = "GDP per capita ($1000s)",
+        y = "Life expectancy (yrs)",
+        caption = "Plot created by @gavin_fay",
+        size = "Population (M)",
+        color = "Continent") +
+   theme_minimal()
```

Recommended reading

An introduction to R (Venables et al.)

– <http://cran.r-project.org/doc/manuals/R-intro.pdf> - Today's material: Chapters 8, 11.

Probability distributions in R

R includes a set of probability distributions that can be used to simulate and model data.

If the function for the probability model is named `xxx`


- ▶ `pxxx`: the cumulative distribution $P(X \leq x)$
- ▶ `dxxx`: the probability distribution/density function $f(x)$
- ▶ `qxxx`: the quantile q , the smallest x such that $P(X \leq x) > q$
- ▶ `rxxx`: generate a random variable from the model `xxx`

Probability distributions in R

Distribution	R name	Additional arguments
beta	<code>beta</code>	shape1, shape2
binomial	<code>binom</code>	size, prob
Cauchy	<code>cauchy</code>	location, scale
chi-squared	<code>chisq</code>	df
exponential	<code>exp</code>	rate
F	<code>f</code>	df1, df2
gamma	<code>gamma</code>	shape, scale
geometric	<code>geom</code>	prob
hypergeometric	<code>hyper</code>	m,n, k
lognormal	<code>lnorm</code>	meanlog, sdlog
logistic	<code>logis</code>	location, scale
negative binomial	<code>nbinom</code>	size, prob
normal	<code>norm</code>	mean,sd
Poisson	<code>pois</code>	lambda
Student's t	<code>t</code>	df
uniform	<code>unif</code>	min, max
Weibull	<code>weibull</code>	shape, scale
Wilcoxon	<code>wilcox</code>	m,n

Standard normal distribution

$$(\mu = 0, \sigma^2 = 1)$$



Functions for normal distribution

Values of x for different quantiles

```
qnorm(p, mean=0, sd=1, lower.tail=TRUE, log.p=FALSE)
```

```
> quants <- qnorm(c(0.01,0.025,0.05,0.95,0.975,0.99))  
> round(quants,2)  
[1] -2.33 -1.96 -1.64  1.64  1.96  2.33
```

$P(X \leq x)$

```
pnorm(q, mean=0, sd=1, lower.tail=TRUE, log.p=FALSE)
```

```
> pnorm(quants)  
[1] 0.010 0.025 0.050 0.950 0.975 0.990
```

Functions for normal distribution

Density (probability 'mass' per unit value of x)

```
> dnorm(quants, mean = 0, sd = 1)
[1] 0.02665214 0.05844507 0.10313564 0.10313564 0.05844507
```

Generating standard normal random variables

```
> rnorm(n=10, mean = 0, sd = 1)
[1] 1.3149588 0.9781675 0.8817912 0.4822047 0.9657529
[7] 0.2839578 -0.1616986 1.9355718 1.7232308
```

Generating random numbers

Often a good idea to use `set.seed()` and save the script detailing which number was used.

This ensures you can exactly repeat your results.

```
> set.seed(42)
> rnorm(3)
[1]  1.3709584 -0.5646982  0.3631284
> rnorm(3)
[1]  0.6328626  0.4042683 -0.1061245
> set.seed(42)
> rnorm(3)
[1]  1.3709584 -0.5646982  0.3631284
> rnorm(3)
[1]  0.6328626  0.4042683 -0.1061245
```

The sample() function

To generate random numbers from discrete sets of values:

- With or without replacement
- Equal or weighted probability

Extremely useful function that underlies many modern statistical techniques:

- Resampling
- Bootstrapping
- Markov-chain Monte-Carlo (MCMC)

e.g. Roll 10 dice

```
> sample(1:6, size = 10, replace = TRUE)
[1] 4 1 5 6 4 2 2 3 1 1
```

Pick 3 species of bear

```
> bears <- c("American Black", "Asian Black", "Brown",
+           "Panda", "Polar", "Sloth", "Spectacled", "Sun")
> sample(bears, size = 3, replace = FALSE)
[1] "Brown" "Panda" "Polar"
```

Lab Exercise 1/3

- a. Generate 1,000 random normal numbers with mean 24 and standard deviation 10. Find the proportion of those random numbers that are ≥ 2 standard deviations from the sample mean.
- b. Flip a (fair) coin six times.
- c. Find the probability of getting six heads on those six flips (i.e. $P(X = 6)$ given $n = 6$).
- d. How much more likely is it to get three heads than six?
- e. For a standard normal random variable, find the number x such that $P(-x \leq X \leq x) = 0.24$.
- f. The mean rate of arrival of alewives at a weir is 3.5 per hour. Plot the probability distribution function for the number of alewife arrivals in an hour.
- g. Find the 95% confidence interval for the number of alewives arriving per day.

Linear models in R

Recall linear regression model:

$$y_i = \sum_{j=0}^p \beta_j x_{ij} + \varepsilon_i, \varepsilon_i \sim N(0, \sigma^2), i = 1, \dots, n$$

In matrix form: $y = X\beta + \varepsilon$

Model formulae in R, $y \sim x$, try ?formula

Formula	Description
$y \sim x1 - 1$	- means leave something out. Fit the slope but not the intercept
$y \sim x1 + x2$	model with covariates $x1$ and $x2$
$y \sim x1 + x2 + x1:x2$	model with covariates $x1$ and $x2$ and an interaction between $x1:x2$
$y \sim x1 * x2$	* denotes factor crossing, and is equivalent to the previous statement
$y \sim (x1 + x2 + x3)^2$	^ indicates crossing to the specified degree. Fit the 3 main effects for $x1$, $x2$, and $x3$ with all possible second order interactions
$y \sim I(x1 + x2)$	I means treat something as is. So the model with single covariate which is the sum of $x1$ and $x2$. (This way we don't have to create the variable $x1 + x2$)

Linear models, `lm()`

The basic function for fitting ordinary multiple models is `lm()`

```
fitted.model <- lm(formula, data = data.frame)
```

e.g. species richness on beaches (Zuur Chapters 5 & 27)

```
> RIKZ <- read_table(file = "../data/RIKZ.txt")  
> RIKZ <- RIKZ |>  
+   mutate(Richness = rowSums(RIKZ[,2:76]>0)) |>  
+   select(Richness, 77:89)  
> RIKZ_lm1 <- lm(Richness ~ NAP, data = RIKZ)
```

Extracting model information

The value of `lm()` is a fitted model object.

- a list of results of class `lm`.

Information about the fitted model can be extracted, displayed, plotted, using some generic functions, including:

- ▶ `anova(object1, object2)` Compares a submodel with an outer model and produces an analysis of variance table
- ▶ `coef(object)` Extract the regression coefficient
- ▶ `deviance(object)` Residual sum of squares
- ▶ `formula(object)` Extract the model formula
- ▶ `plot(object)` Produce four plots, showing residuals, fitted values and some diagnostics
- ▶ `predict(object, newdata=data.frame)` Model predictions on new data
- ▶ `residuals(object)` or `resid(object)` Extract the matrix of residuals
- ▶ `step(object)` forward or backward model selection using AIC
- ▶ `summary(object)` Print a comprehensive summary of the results
- ▶ `vcov(object)` Return variance-covariance matrix of main parameters

Summary objects

```
summary(lm1)
```

```
Call:
```

```
lm(formula = Richness ~ NAP, data = RIKZ)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-5.0675	-2.7607	-0.8029	1.3534	13.8723

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	6.6857	0.6578	10.164	5.25e-13	***
NAP	-2.8669	0.6307	-4.545	4.42e-05	***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.16 on 43 degrees of freedom
```

```
Multiple R-squared:  0.3245,    Adjusted R-squared:  0.3088
```

```
F-statistic: 20.66 on 1 and 43 DF,  p-value: 4.418e-05
```

broom helper functions

```
broom::tidy()
```

```
> library(broom)
> tidy(RIKZ_lm1, conf.int = TRUE)
# A tibble: 2 x 7
```

	term	estimate	std.error	statistic	p.value	conf.low	conf.high
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	6.69	0.658	10.2	5.25e-13	5.36	8.02
2	NAP	-2.87	0.631	-4.55	4.42e- 5	-4.14	-1.60

```
broom::glance()
```

```
> glance(RIKZ_lm1)
# A tibble: 1 x 12
```

	r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.325	0.309	4.16	20.7	0.0000442	1	-127.

```
# ... with 3 more variables: deviance <dbl>, df.residual <int>,
# AICc <dbl>
```

{modelsummary}

The {modelsummary} package provides a wealth of options for creating publication-quality summary tables from fitted model objects, with bindings for most of the types of models we are using in this class.

There are lots of options....

```
> library(modelsummary)
> modelsummary(list(lm1 = RIKZ_lm1),
+               statistic = "conf.int",
+               stars = TRUE)
```

lm1	
(Intercept)	6.686***
	[5.359, 8.012]
NAP	-2.867***
	[-4.139, -1.595]
Num.Obs.	45
R2	0.325
R2 Adj.	0.309
AIC	260.0
BIC	265.4
Log.Lik.	-126.977
F	20.660
RMSE	4.07

Are model assumptions met?

In the MASS library there are many functions.

Are samples independent? (Sample design.)

Residuals normally distributed?

- Histograms, qq-plots: `qqplot()` and `qqline()`
- Kolmogorov-Smirnov normality test: `ks.test()`
- Shapiro-Wilk normality test: `shapiro.test()`

Similar variance among samples?

- Boxplots
- Bartlett's test for equal variance: `bartlett.test()`
- Fligner-Killeen test for equal variance: `fligner.test()`


Checking assumptions

Model assumptions can be evaluated by plotting the model object.


```
> plot(RIKZ_lm1)
> # OR
> library(ggglm)
> ggglm(RIKZ_lm1)
```

```
> gg1m(RIKZ_lm1)
```


Residuals vs Fitted




Normal Q-Q



Scale-Location



Residual vs. Leverage



Extracting model fit & residuals

```
broom::augment()
```

```
> RIKZ_lm1_aug <- augment(RIKZ_lm1, se_fit = )  
> glimpse(RIKZ_lm1_aug)
```

```
Rows: 45
```

```
Columns: 8
```

```
$ Richness    <dbl> 11, 10, 13, 11, 10, 8, 9, 8, 19, 17, 6,  
$ NAP         <dbl> 0.045, -1.036, -1.336, 0.616, -0.684, 1,  
$ .fitted     <dbl> 6.556653, 9.655722, 10.515778, 4.919680,  
$ .resid      <dbl> 4.4433465, 0.3442783, 2.4842223, 6.08031,  
$ .hat        <dbl> 0.02432839, 0.06623475, 0.08738853, 0.02,  
$ .sigma      <dbl> 4.151533, 4.208801, 4.189991, 4.100641,  
$ .cooksds    <dbl> 0.0145788938, 0.0002601525, 0.0187094818,  
$ .std.resid  <dbl> 1.08136537, 0.08564557, 0.62511744, 1.47
```


Extracting model fit & residuals


```
> RIKZ_lm1_aug <- augment(RIKZ_lm1,  
+                           se_fit = TRUE)  
> glimpse(RIKZ_lm1_aug)  
Rows: 45  
Columns: 9  
$ Richness    <dbl> 11, 10, 13, 11, 10, 8, 9, 8, 19, 17, 6,  
$ NAP         <dbl> 0.045, -1.036, -1.336, 0.616, -0.684, 1,  
$ .fitted     <dbl> 6.556653, 9.655722, 10.515778, 4.919680,  
$ .se.fit     <dbl> 0.6488474, 1.0706040, 1.2297395, 0.64280,  
$ .resid      <dbl> 4.4433465, 0.3442783, 2.4842223, 6.08031,  
$ .hat        <dbl> 0.02432839, 0.06623475, 0.08738853, 0.02,  
$ .sigma      <dbl> 4.151533, 4.208801, 4.189991, 4.100641,  
$ .cooksdi    <dbl> 0.0145788938, 0.0002601525, 0.0187094818,  
$ .std.resid  <dbl> 1.08136537, 0.08564557, 0.62511744, 1.47
```

Extracting model fit & residuals

```
> RIKZ_lm1_aug <- augment(RIKZ_lm1,  
+                           interval = "confidence")  
> glimpse(RIKZ_lm1_aug)  
Rows: 45  
Columns: 10  
$ Richness    <dbl> 11, 10, 13, 11, 10, 8, 9, 8, 19, 17, 6,  
$ NAP         <dbl> 0.045, -1.036, -1.336, 0.616, -0.684, 1,  
$ .fitted     <dbl> 6.556653, 9.655722, 10.515778, 4.919680,  
$ .lower      <dbl> 5.24812804, 7.49664302, 8.03577164, 3.62,  
$ .upper      <dbl> 7.865179, 11.814800, 12.995784, 6.216014,  
$ .resid      <dbl> 4.4433465, 0.3442783, 2.4842223, 6.08031,  
$ .hat        <dbl> 0.02432839, 0.06623475, 0.08738853, 0.02,  
$ .sigma      <dbl> 4.151533, 4.208801, 4.189991, 4.100641,  
$ .cooksd     <dbl> 0.0145788938, 0.0002601525, 0.0187094818,  
$ .std.resid  <dbl> 1.08136537, 0.08564557, 0.62511744, 1.47
```


Extracting model fit & residuals

```
> RIKZ_lm1_aug |>  
+   janitor::clean_names() |>  
+   ggplot() +  
+   aes(x = nap, y = fitted) +  
+   geom_line() +  
+   geom_point(aes(x = nap, y = richness))
```




Fitting models directly on plots

```
> ggplot(RIKZ, aes(x = NAP, y = Richness)) +  
+   geom_point() +  
+   geom_smooth(method = "lm")  
`geom_smooth()` using formula 'y ~ x'
```



Building and Comparing models

```
> library(moderndiver)
> ggplot(RIKZ) +
+   aes(x = NAP, y = Richness, color = factor(week)) +
+   geom_point() +
+   labs(x = "NAP", y = "Species Richness",
+         color = "Week") +
+   geom_parallel_slopes(se = TRUE)
```



Building and Comparing models

An alternative model for the RIKZ species richness is:

```
> RIKZ_lm2 <- lm(Richness ~ NAP + factor(week), data = RIKZ)
> #Can also create this model using the `update()` function:
> RIKZ_lm2 <- update(RIKZ_lm1, .~. + factor(week))
```

Compare models with AIC() and via anova()

```
> anova(RIKZ_lm1, RIKZ_lm2)
```

Analysis of Variance Table

Model 1: Richness ~ NAP

Model 2: Richness ~ NAP + factor(week)

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	43	744.12				
2	40	357.00	3	387.11	14.458	1.581e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Lab exercise 2/3

1. Extract the residuals from the RIKZ_lm2 model
2. Are the linear regression assumptions met? Explain your reasoning
3. Summarize the results of the model. What are the parameter estimates telling you about species richness on these beaches?

Model predictions for new data

Both `augment()` and the `predict()` function can be used to obtain predictions from a fitted model object to a new data frame.

This can be useful when:

- a subset of the data was reserved from the fitting process ('test data')
- you want to obtain model predictions at values other than the data (for example when plotting fitted values)

```
> newdata <- RIKZ |>
+   slice(1:10) |>
+   mutate(NAP = NAP + 1)
> new_predictions <- augment(RIKZ_lm1, newdata = newdata)
```


More on interrogating model results

You can also interrogate model objects directly. Type `names(object)` to get a list of the components of object.

```
> names(RIKZ_lm1)
[1] "coefficients" "residuals"      "effects"         "rank"
[5] "fitted.values" "assign"          "qr"              "df.residual"
[9] "xlevels"       "call"           "terms"           "model"
```

`str()` can also be used.

More on interrogating model results

Note that `summary(object)` is also a list, and components can be extracted from here too.

```
> names(summary(RIKZ_lm1))
[1] "call"          "terms"          "residuals"      "coefficients"
[5] "aliased"        "sigma"          "df"             "r.squared"
[9] "adj.r.squared" "fstatistic"     "cov.unscaled"
> str(summary(RIKZ_lm1))
List of 11
 $ call      : language lm(formula = Richness ~ NAP, data = dat)
 $ terms     :Classes 'terms', 'formula' language Richness ~ NAP
 .. ..- attr(*, "variables")= language list(Richness, NAP)
 .. ..- attr(*, "factors")= int [1:2, 1] 0 1
 .. ..- attr(*, "dimnames")=List of 2
 .. .. $ : chr [1:2] "Richness" "NAP"
 .. .. $ : chr "NAP"
 .. ..- attr(*, "term.labels")= chr "NAP"
 .. ..- attr(*, "order")= int 1
 .. ..- attr(*, "intercept")= int 1
```

Centering covariates

When dealing with numeric covariates, it often makes sense to center these variables before performing a regression.

- i.e. subtract the mean from each covariate.


This helps to make our parameters more meaningful.

- they now correspond to the average case in the data, rather than some extrapolated value.
- this is particularly useful for intercept parameters.

Fitting multiple models simultaneously

We may be interested in fitting many (many) models.
R has functionality to do this efficiently.


```
> library(broom)
> ggplot(gapminder, aes(x=year, y=lifeExp, group = country)) +
+   geom_line() +
+   facet_wrap(~continent)
```




Fitting multiple models simultaneously

```
> gapminder_models <- gapminder |>
+   group_by(country) |>
+   nest() |>
+   mutate(model = map(data, ~lm(lifeExp~year, data = .x))) |>
+   mutate(coefs = map(model, tidy, conf.int = TRUE)) |>
+   unnest(coefs)
```

```
> ggplot(gapminder_models, aes(x = estimate, group = term)) +  
+   geom_histogram(fill="black", col="white") +  
+   facet_wrap(~term, scales = "free")
```



```
> slopes <- filter(gapminder_models, term == "year") %>%  
+   arrange(desc(estimate))  
> continents <- select(gapminder, country, continent) |>  
+   distinct()  
> slopes <- slopes |> left_join(continents)  
Joining, by = "country"  
> ggplot(slopes, aes(x = fct_reorder(country, estimate), y = estimate)) +  
+   #geom_histogram(fill="black", col="white") +  
+   geom_point(alpha=0.5) +  
+   geom_errorbar(aes(ymin=conf.low, ymax=conf.high), width=.05) +  
+   coord_flip() +  
+   labs(y = "slope",  
+        x = "Country") +  
+   theme_minimal()
```



Lab exercise 3/3

1. Fit a linear regression using 2007 gapminder data of the form `lm(gdpPercap ~ continent, data = gapminder)`, where `gdpPercap` is the new outcome variable `y`. Get information about the best-fitting line from the regression table. How do the regression results match up with those of an analysis of life expectancy by continent?
2. Extract the model coefficients and their 95 percent confidence intervals.
3. Plot the residuals vs the fitted values and comment on their distribution and patterns.
4. Identify the five countries with the five most negative residuals? What do these negative residuals say about their life expectancy relative to their continents life expectancy?
5. Repeat this process, but identify the five countries with the five most positive residuals. What do these positive residuals say about their life expectancy relative to their continents life expectancy?

